

# From Scratch Implementation of Deep Reinforcement Learning Algorithms for Continuous Control

Ruturaj Sambhus

Virginia Tech

December 1, 2022

## 1 Theory

- Policy Gradient Algorithms - REINFORCE
- Actor Critic Algorithms - A2C

## 2 Implementation Details

- Framework
- Neural Networks

## 3 Results

- Gym and PyBullet Results
- Sensitivity to Hyper-parameters

# Class of Reinforcement Learning Methods

## Value Based Methods

- The value function (expected return) of the state space is learned
- This is used for discrete state and action spaces
- Algorithms - Deep Q Learning (DQN), Dueling DQN

## Policy Gradients

- The policy (action given state) for state space is learned
- Can be used for continuous state and action spaces
- Algorithms - REINFORCE or Vanilla Policy Gradient, TRPO, PPO

## Actor Critic

- Combination of both value-based and policy gradient methods
- Can be used for continuous state and action spaces
- Algorithms - Advantage Actor Critic (A2C), A3C, DDPG, SAC

# REINFORCE - Vanilla Policy Gradient

## REINFORCE with Baseline (episodic), for estimating $\pi_{\theta} \approx \pi_*$

Input: a differentiable policy parameterization  $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization  $\hat{v}(s, \mathbf{w})$

Algorithm parameters: step sizes  $\alpha^{\theta} > 0$ ,  $\alpha^{\mathbf{w}} > 0$

Initialize policy parameter  $\theta \in \mathbb{R}^{d'}$  and state-value weights  $\mathbf{w} \in \mathbb{R}^d$  (e.g., to  $\mathbf{0}$ )

Loop forever (for each episode):

    Generate an episode  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ , following  $\pi(\cdot|\cdot, \theta)$

    Loop for each step of the episode  $t = 0, 1, \dots, T-1$ :

$$G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k \quad (G_t)$$

$$\delta \leftarrow G - \hat{v}(S_t, \mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha^{\mathbf{w}} \delta \nabla \hat{v}(S_t, \mathbf{w})$$

$$\theta \leftarrow \theta + \alpha^{\theta} \gamma^t \delta \nabla \ln \pi(A_t|S_t, \theta)$$

Figure: REINFORCE with Baseline [Sutton and Barto, 2018]

# A2C - Advantage Actor Critic

## One-step Actor-Critic (episodic), for estimating $\pi_{\theta} \approx \pi_*$

Input: a differentiable policy parameterization  $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization  $\hat{v}(s, \mathbf{w})$

Parameters: step sizes  $\alpha^{\theta} > 0$ ,  $\alpha^{\mathbf{w}} > 0$

Initialize policy parameter  $\theta \in \mathbb{R}^{d'}$  and state-value weights  $\mathbf{w} \in \mathbb{R}^d$  (e.g., to  $\mathbf{0}$ )

Loop forever (for each episode):

    Initialize  $S$  (first state of episode)

$I \leftarrow 1$

    Loop while  $S$  is not terminal (for each time step):

$A \sim \pi(\cdot|S, \theta)$

        Take action  $A$ , observe  $S', R$

$\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$       (if  $S'$  is terminal, then  $\hat{v}(S', \mathbf{w}) \doteq 0$ )

$\mathbf{w} \leftarrow \mathbf{w} + \alpha^{\mathbf{w}} \delta \nabla \hat{v}(S, \mathbf{w})$

$\theta \leftarrow \theta + \alpha^{\theta} I \delta \nabla \ln \pi(A|S, \theta)$

$I \leftarrow \gamma I$

$S \leftarrow S'$

Figure: One Step A2C [Sutton and Barto, 2018]

# Way Things Move




Figure: Framework

# Class Based Structure

- Many moving parts and hence divided into following classes
  - Networks - policy and value neural networks
  - Memory - to store batch data generated during policy evaluation
  - Agent - performs policy evaluation and policy update
- Agent has memory, actor, and critic networks
- Agent plays a given policy for a specified number of time steps and collects  $\{s, a, r, s', terminal\}$  in the memory
- Reinforce - rewards are converted into returns and stored
- Agent performs policy update
- Advantage (TD-error) is calculated for the batch number of states and backward pass is performed

# Neural Network

- Choosing PyTorch - easy to install, less learning curve to use
- Policy outputs mean and variance of Normal distribution for each action dimension
- Multiple ways to do this - some work and some do not
- I found a working implementation with the way shown below

## Example (Definition)

```
self.fc1 = nn.Linear(*self.input_dim, self.fc1_size)
self.fc2 = nn.Linear(self.fc1_size, self.fc2_size)
self.mu = nn.Linear(self.fc2_size, self.action_dim)
logstds_param = nn.Parameter \
    (torch.full((self.action_dim,), 0.1))
self.register_parameter('logstds', logstds_param)
sigma = torch.clamp(self.logstds.exp(), min=1e-3, max=50)
```







# Hyperparameters





Parameters	Values
Policy-network	$\{n(\text{state}), 64, 64, n(\text{actions})\}$
Value-network	$\{n(\text{state}), 64, 64, 1\}$
Activation	ReLU
Actor-lr	0.0005
Critic-lr	0.005
Discount $\gamma$	0.95
Evaluation Batch (timesteps)	64 (REINFORCE), 32 (A2C)
Gradient clip norm	0.5
Policy Standard Deviation Limits	min=1e-3, max = 50

Table: Parameters for Gym Pendulum-v1 Environment





# A2C Results - Pendulum-v1







# A2C Results - Other Environments



# REINFORCE Results - Pendulum-v1



# REINFORCE Results - Other Environments



# How Important Is Gradient Clipping?

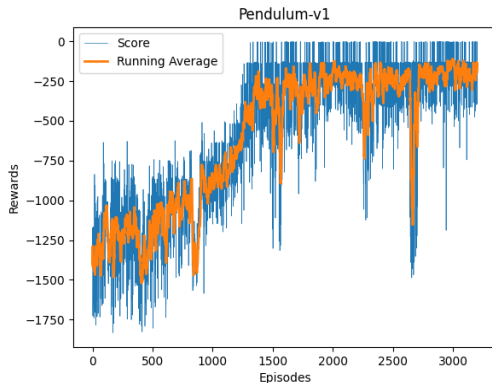


Figure: REINFORCE Gradient Clipping =  $1e12$  for Pendulum-v1

Gradient Clipping does not affect performance!

# How Important Is Gradient Clipping?





Figure: A2C Gradient Clipping = 100 and 10 for Pendulum-v1

Gradient Clipping must!

# Effect of Evaluation Batch Size

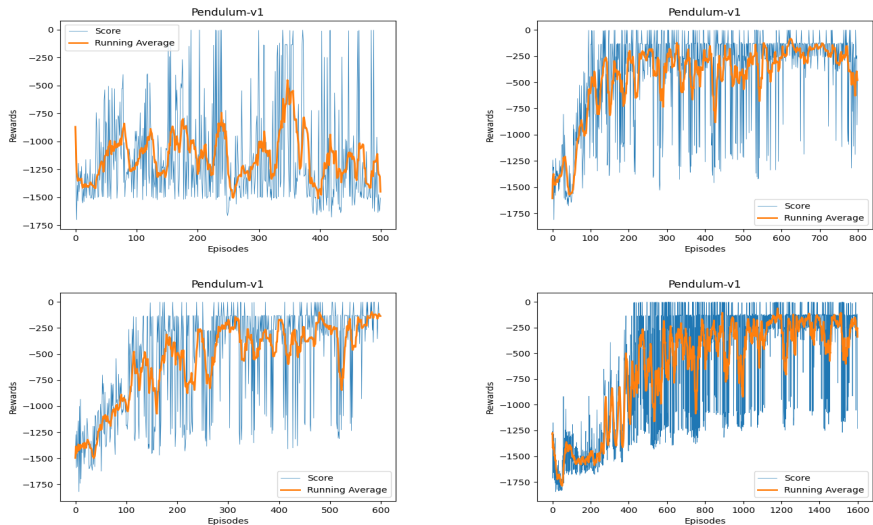


Figure: A2C evaluation batch = 1, 8, 16 and 64 for Pendulum-v1



# Effect of Evaluation Batch Size

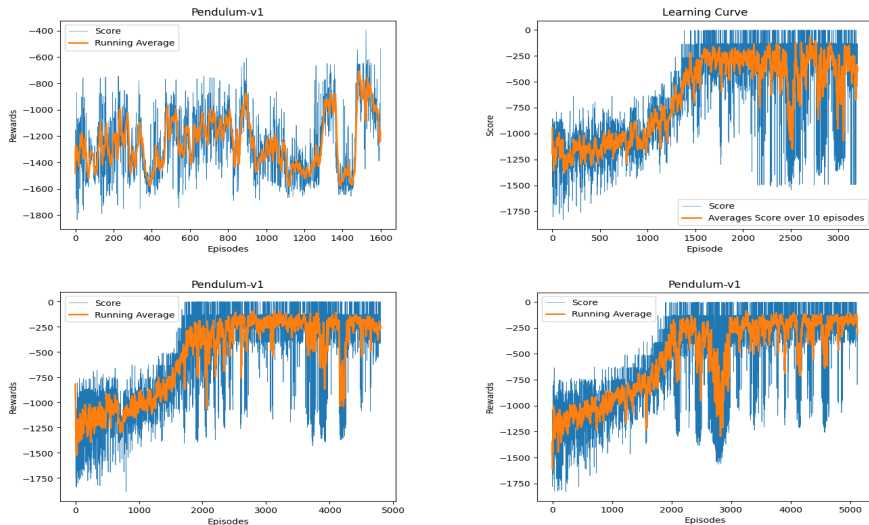


Figure: REINFORCE evaluation batch = 32, 64, 128 and 256 for Pendulum-v1

# Effect of Neural Net Sizes







Figure: Neurons per layer 32, 128 (A2C and REINFORCE) Pendulum-v1

# Observations and Discussions

- REINFORCE needs more samples than A2C agreeing with the theory
- Both the algorithms fail to learn Mountain Car Continuous, a sparse reward environment
- Need advanced algorithms for a high degree of freedom control tasks
- A2C doesn't work without gradient clipping
- A2C has an optimal sampling batch size, single step A2C did not work; REINFORCE needs a minimum batch size after which it does not affect performance much
- Increased sizes of Neural Nets may not give better performance
- Deep RL needs a lot of hyper-parameter tuning!

# References



Sutton, Richard S., and Andrew G. Barto.

Reinforcement learning: An introduction

*Journal Name* MIT press, 2018



Phil Tabor

<https://www.youtube.com/@MachineLearningwithPhil>

*Github* <https://github.com/philtabor/Youtube-Code-Repository>



hermesdt <https://github.com/hermesdt>

*Github* <https://github.com/hermesdt/reinforcement-learning/tree/master/a2c>

# Thank you!